

KTrussExplorer: Exploring the Design Space of K-truss Decomposition Optimizations on GPUs

Safaa Diab, Mhd Ghaith Olabi, Izzat El Hajj

American University of Beirut

HPEC Graph Challenge

September 23, 2020




Overview

KTrussExplorer is a highly parameterized framework for exploring different combinations of k-truss decomposition optimizations on GPUs

Supported features:

- Edge-centric parallelization
- Undirected or directed graphs
 - Directed by index or by degree
- Tiling the adjacency matrix
- Parallelizing intersections
- Removing or marking weak edges
- Recomputing for all or affected edges


Contributions:

- A survey of optimizations
 - A framework for exploring the design space
-  github.com/ielhaji/ktruss-explorer
- A view of the design space
 - Unexplored combinations faster than prior champions


Methodology

- **Software:** KtrussExplorer kernels are implemented in CUDA
- **System:** Evaluation is on one Volta V100 GPU with 16GB of memory
- **Datasets:** We evaluate with all graphs in the graph challenge collection
 - Except: Friendster, graph500-scale24-ef16, and graph500-scale25-ef16 due to limited device memory capacity.
- **Search space:** Design space is searched exhaustively
 - Except: very large graphs

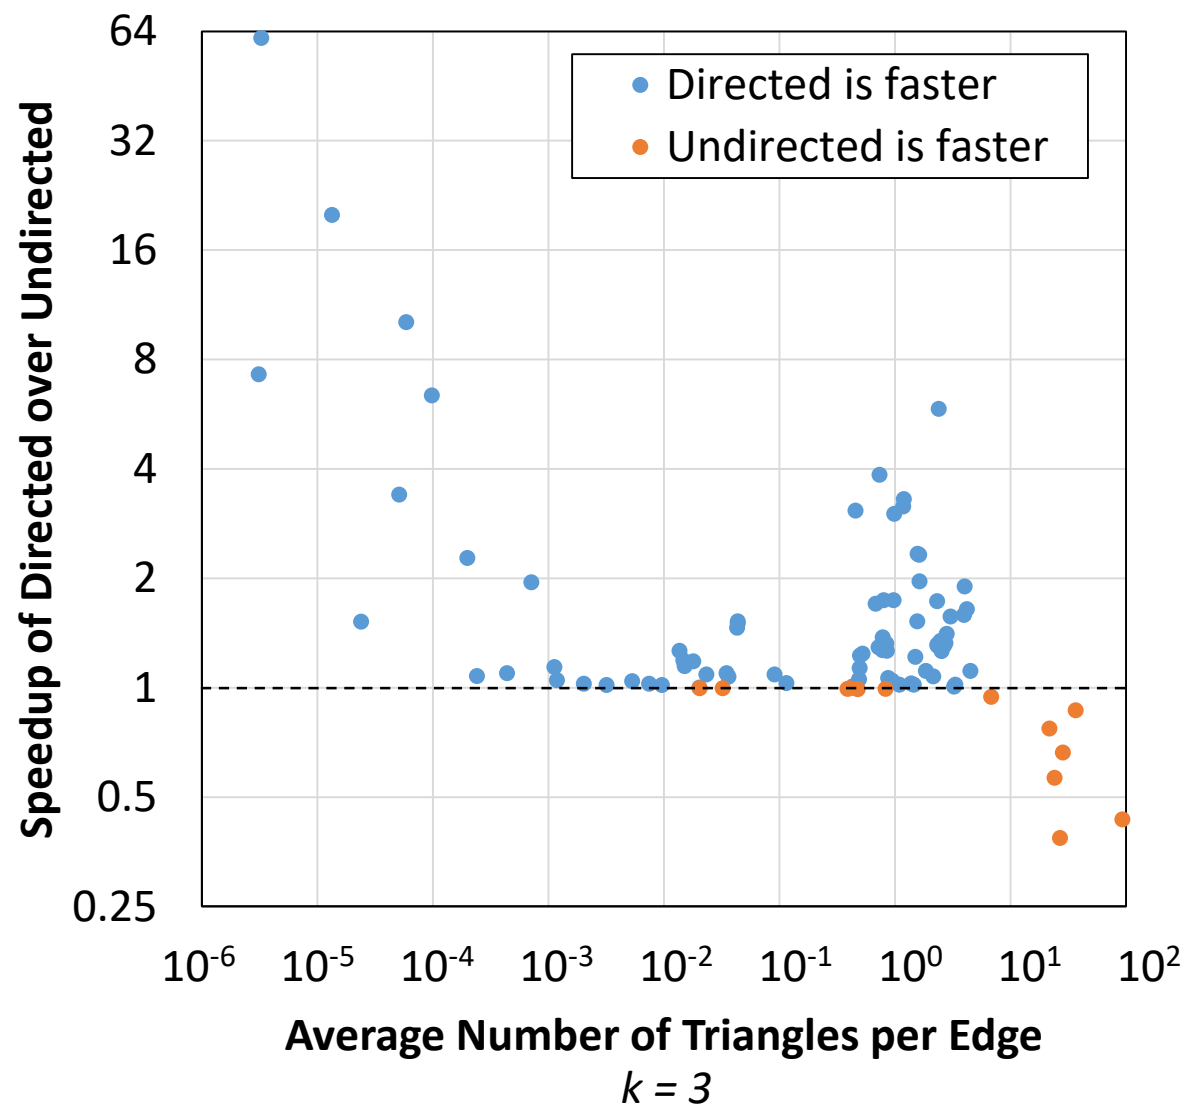
Graph Directedness



- ✓ Less synchronization (no atomics)
- ✓ Stop counting early



- ✓ Less redundancy



Directing Edges by Degree

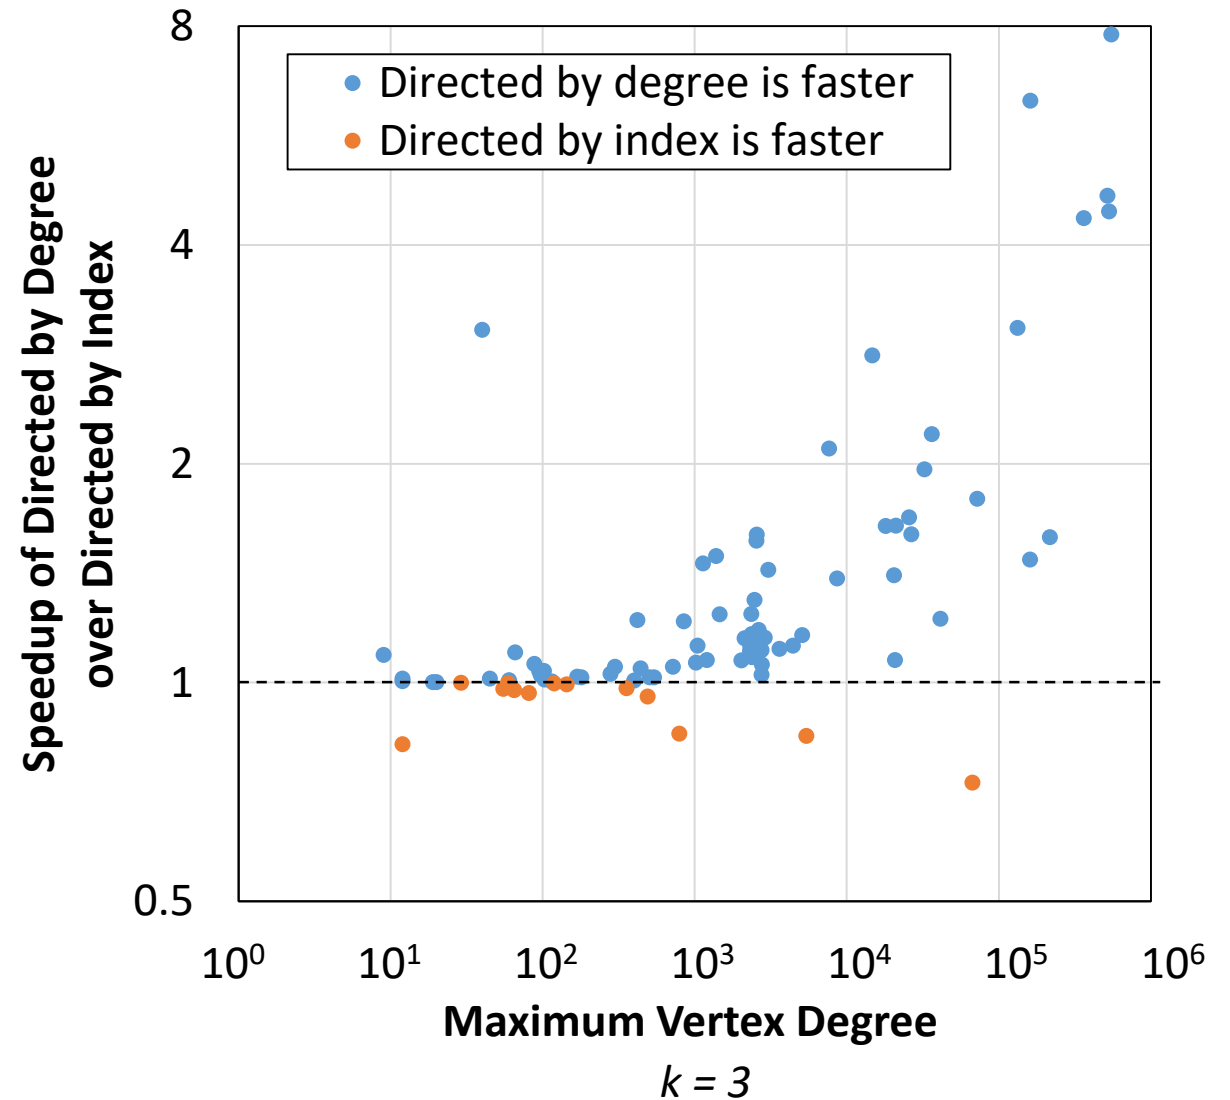
Directed by index


- Keep edges from vertex with lower index to vertex with higher index

Directed by degree

- Keep edges from vertex with lower degree to vertex with higher degree

✓ Advantage: shrink large adjacency lists to reduce load imbalance





Example Graph

	0	1	2	3	4	5	6	7
0		1				1	1	1
1	1			1		1		
2					1	1		1
3		1				1		
4			1					1
5	1	1	1	1				
6	1							1
7	1		1		1		1	

Logical Adjacency List
without Tiling

	0	1	2	3	4	5	6	7
0		1				1	1	1
1	1			1		1		
2					1	1		1
3		1				1		
4			1					1
5	1	1	1	1				
6	1							1
7	1		1		1		1	

Logical Adjacency List
with Tiling

Tiling

srcPtr

0	4	7	10	12	14	18	29	24
---	---	---	----	----	----	----	----	----

dstIdx

1	5	6	7	0	3	5	4	5	7	1	5	2	7	0	1	2	3	0	7	0	2	4	6
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

CSR Representation

✓ Better locality

✓ Partitioning intersections into smaller sub-intersections

srcPtr

0	1	3	3	4	7	8	11	12	13	17	18	20	21	21	22	24
---	---	---	---	---	---	---	----	----	----	----	----	----	----	----	----	----

dstIdx

1	0	3	1	5	6	7	5	4	5	7	5	2	0	1	2	3	0	0	2	7	7	4	6
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Tiled CSR Representation


Benefits of Tiling

- Bad locality




✓ Good locality

✓ Good locality



✓ Good locality



Memory Access Pattern without Tiling




Memory Access Pattern with Tiling








Benefits of Tiling

Intersection



Sub-intersection 1
(trivially empty)


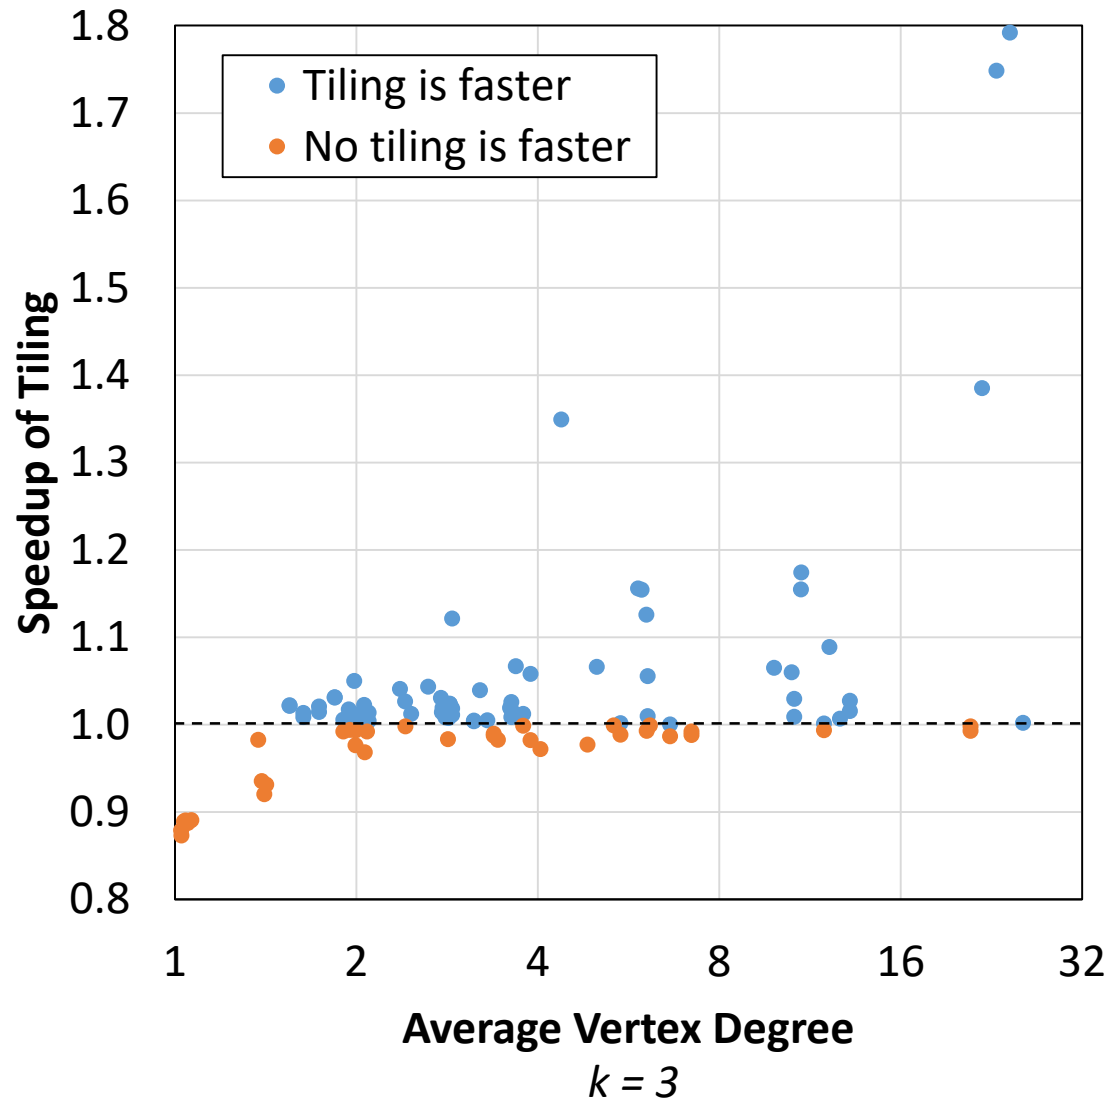
Sub-intersection 2
(trivially empty)




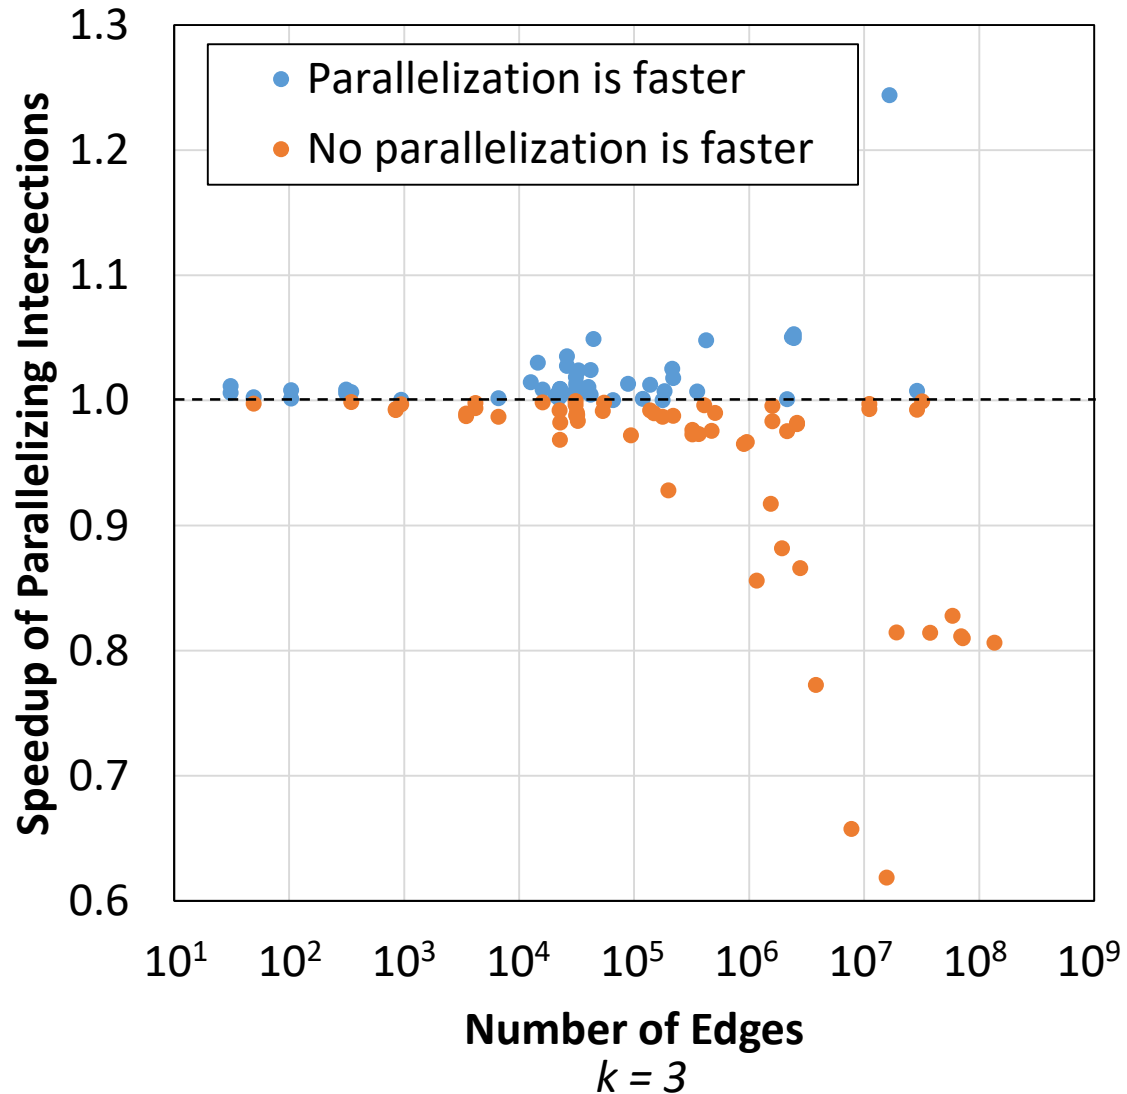
Intersection without Tiling

Intersection with Tiling


Benefits of Tiling




Parallelizing Intersections



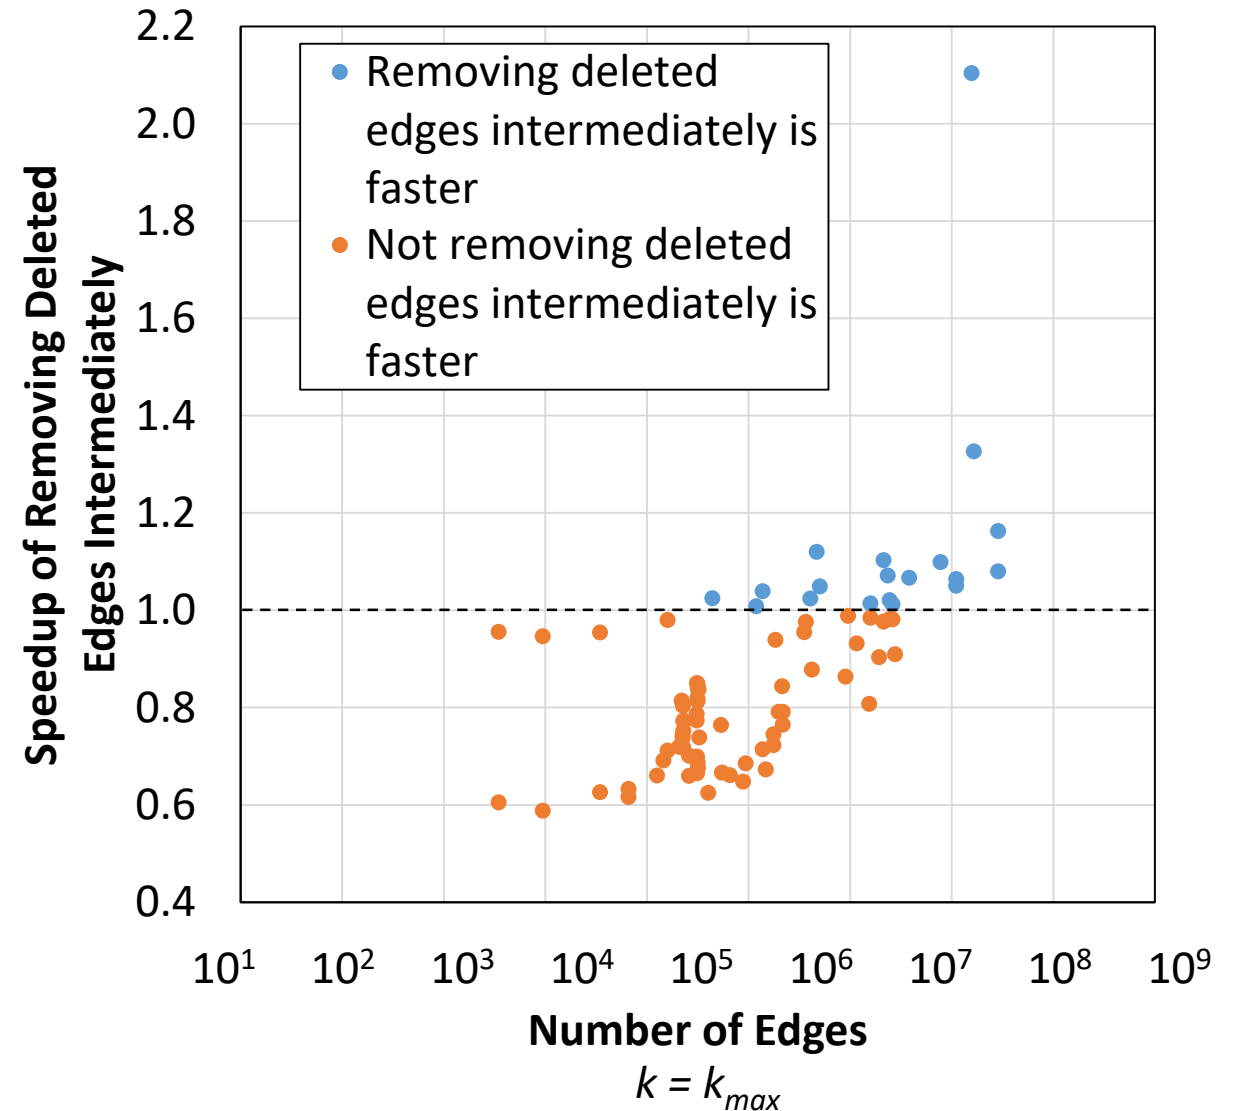

Removing Deleted Edges Intermediately



Mark deleted edges




Remove deleted edges (for select iterations)




Recomputing Support for All or Affected Edges

- Edges that are not affected and whose threads do not need to recount
- Edges that are not affected but whose threads need to recount on behalf of affected edges
- .-.-→ Edges that are affected and whose threads need to recount
-→ Weak edges that were deleted



Undirected Graph



Directed Graph

Graphs performing better with only affected edges reprocessed:


- graph500-scale20-ef16
- graph500-scale21-ef16
- graph500-scale23-ef16

For further investigation:

- Recomputing for affected edges on select iterations (later iterations)

Marking Affected Edges

- Edges that are not affected and whose threads do not need to recount
- Edges that are not affected but whose threads need to recount on behalf of affected edges
- .-.-.-→ Edges that are affected and whose threads need to recount
-→ Weak edges that were deleted




```
01: parallel for  $e = \{u, v\} \in E$  do
02:   if  $e$  is deleted then
03:     mark  $u$  as affected, mark  $v$  as affected
```

Pseudocode for Marking Affected Edges

Marking Affected Edges

- Edges that are not affected and whose threads do not need to recount
- Edges that are not affected but whose threads need to recount on behalf of affected edges
- .-.-.-.-→ Edges that are affected and whose threads need to recount
-→ Weak edges that were deleted




```
01: parallel for  $e = \{u, v\} \in E$  do
02:   if  $e$  is deleted then
03:     mark  $u$  as affected, mark  $v$  as affected
04:   parallel for  $e = \{u, v\} \in E$  do
05:     if  $e$  is not deleted and ( $u$  is affected or  $v$  is affected) then
06:       mark  $e$  as affected
07:       if  $u$  is not affected then mark  $u$  as needs to recount
08:       else if  $v$  is not affected then mark  $v$  as needs to recount
```

Pseudocode for Marking Affected Edges

Marking Affected Edges

- Edges that are not affected and whose threads do not need to recount
- Edges that are not affected but whose threads need to recount on behalf of affected edges
- .-.-.-.-→ Edges that are affected and whose threads need to recount
-→ Weak edges that were deleted

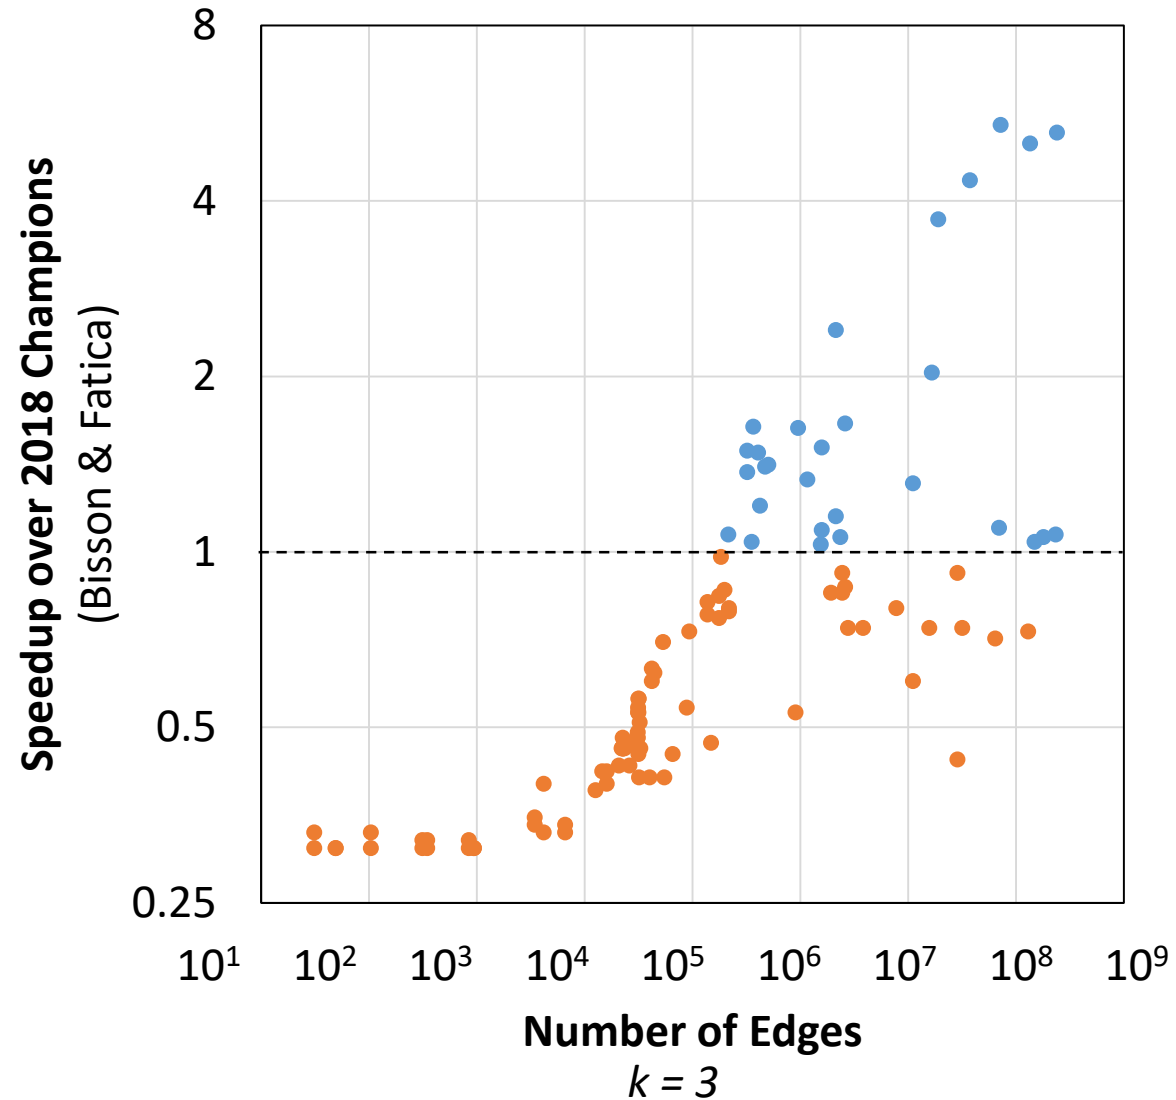


```

01: parallel for  $e = \{u, v\} \in E$  do
02:   if  $e$  is deleted then
03:     mark  $u$  as affected, mark  $v$  as affected
04:   parallel for  $e = \{u, v\} \in E$  do
05:     if  $e$  is not deleted and ( $u$  is affected or  $v$  is affected) then
06:       mark  $e$  as affected
07:       if  $u$  is not affected then mark  $u$  as needs to recount
08:       else if  $v$  is not affected then mark  $v$  as needs to recount
09:   parallel for  $e = \{u, v\} \in E$  do
10:     if  $e$  is not deleted and  $e$  is not affected then
11:       if  $u$  needs to recount or  $v$  needs to recount then
12:         mark  $e$  as needs to recount
  
```

Pseudocode for Marking Affected Edges

Comparison with Prior Champions



KTrussExplorer: Exploring the Design Space of K-truss Decomposition Optimizations on GPUs

Safaa Diab, Mhd Ghaith Olabi, Izzat El Hajj

American University of Beirut



github.com/ielhaji/ktruss-explorer